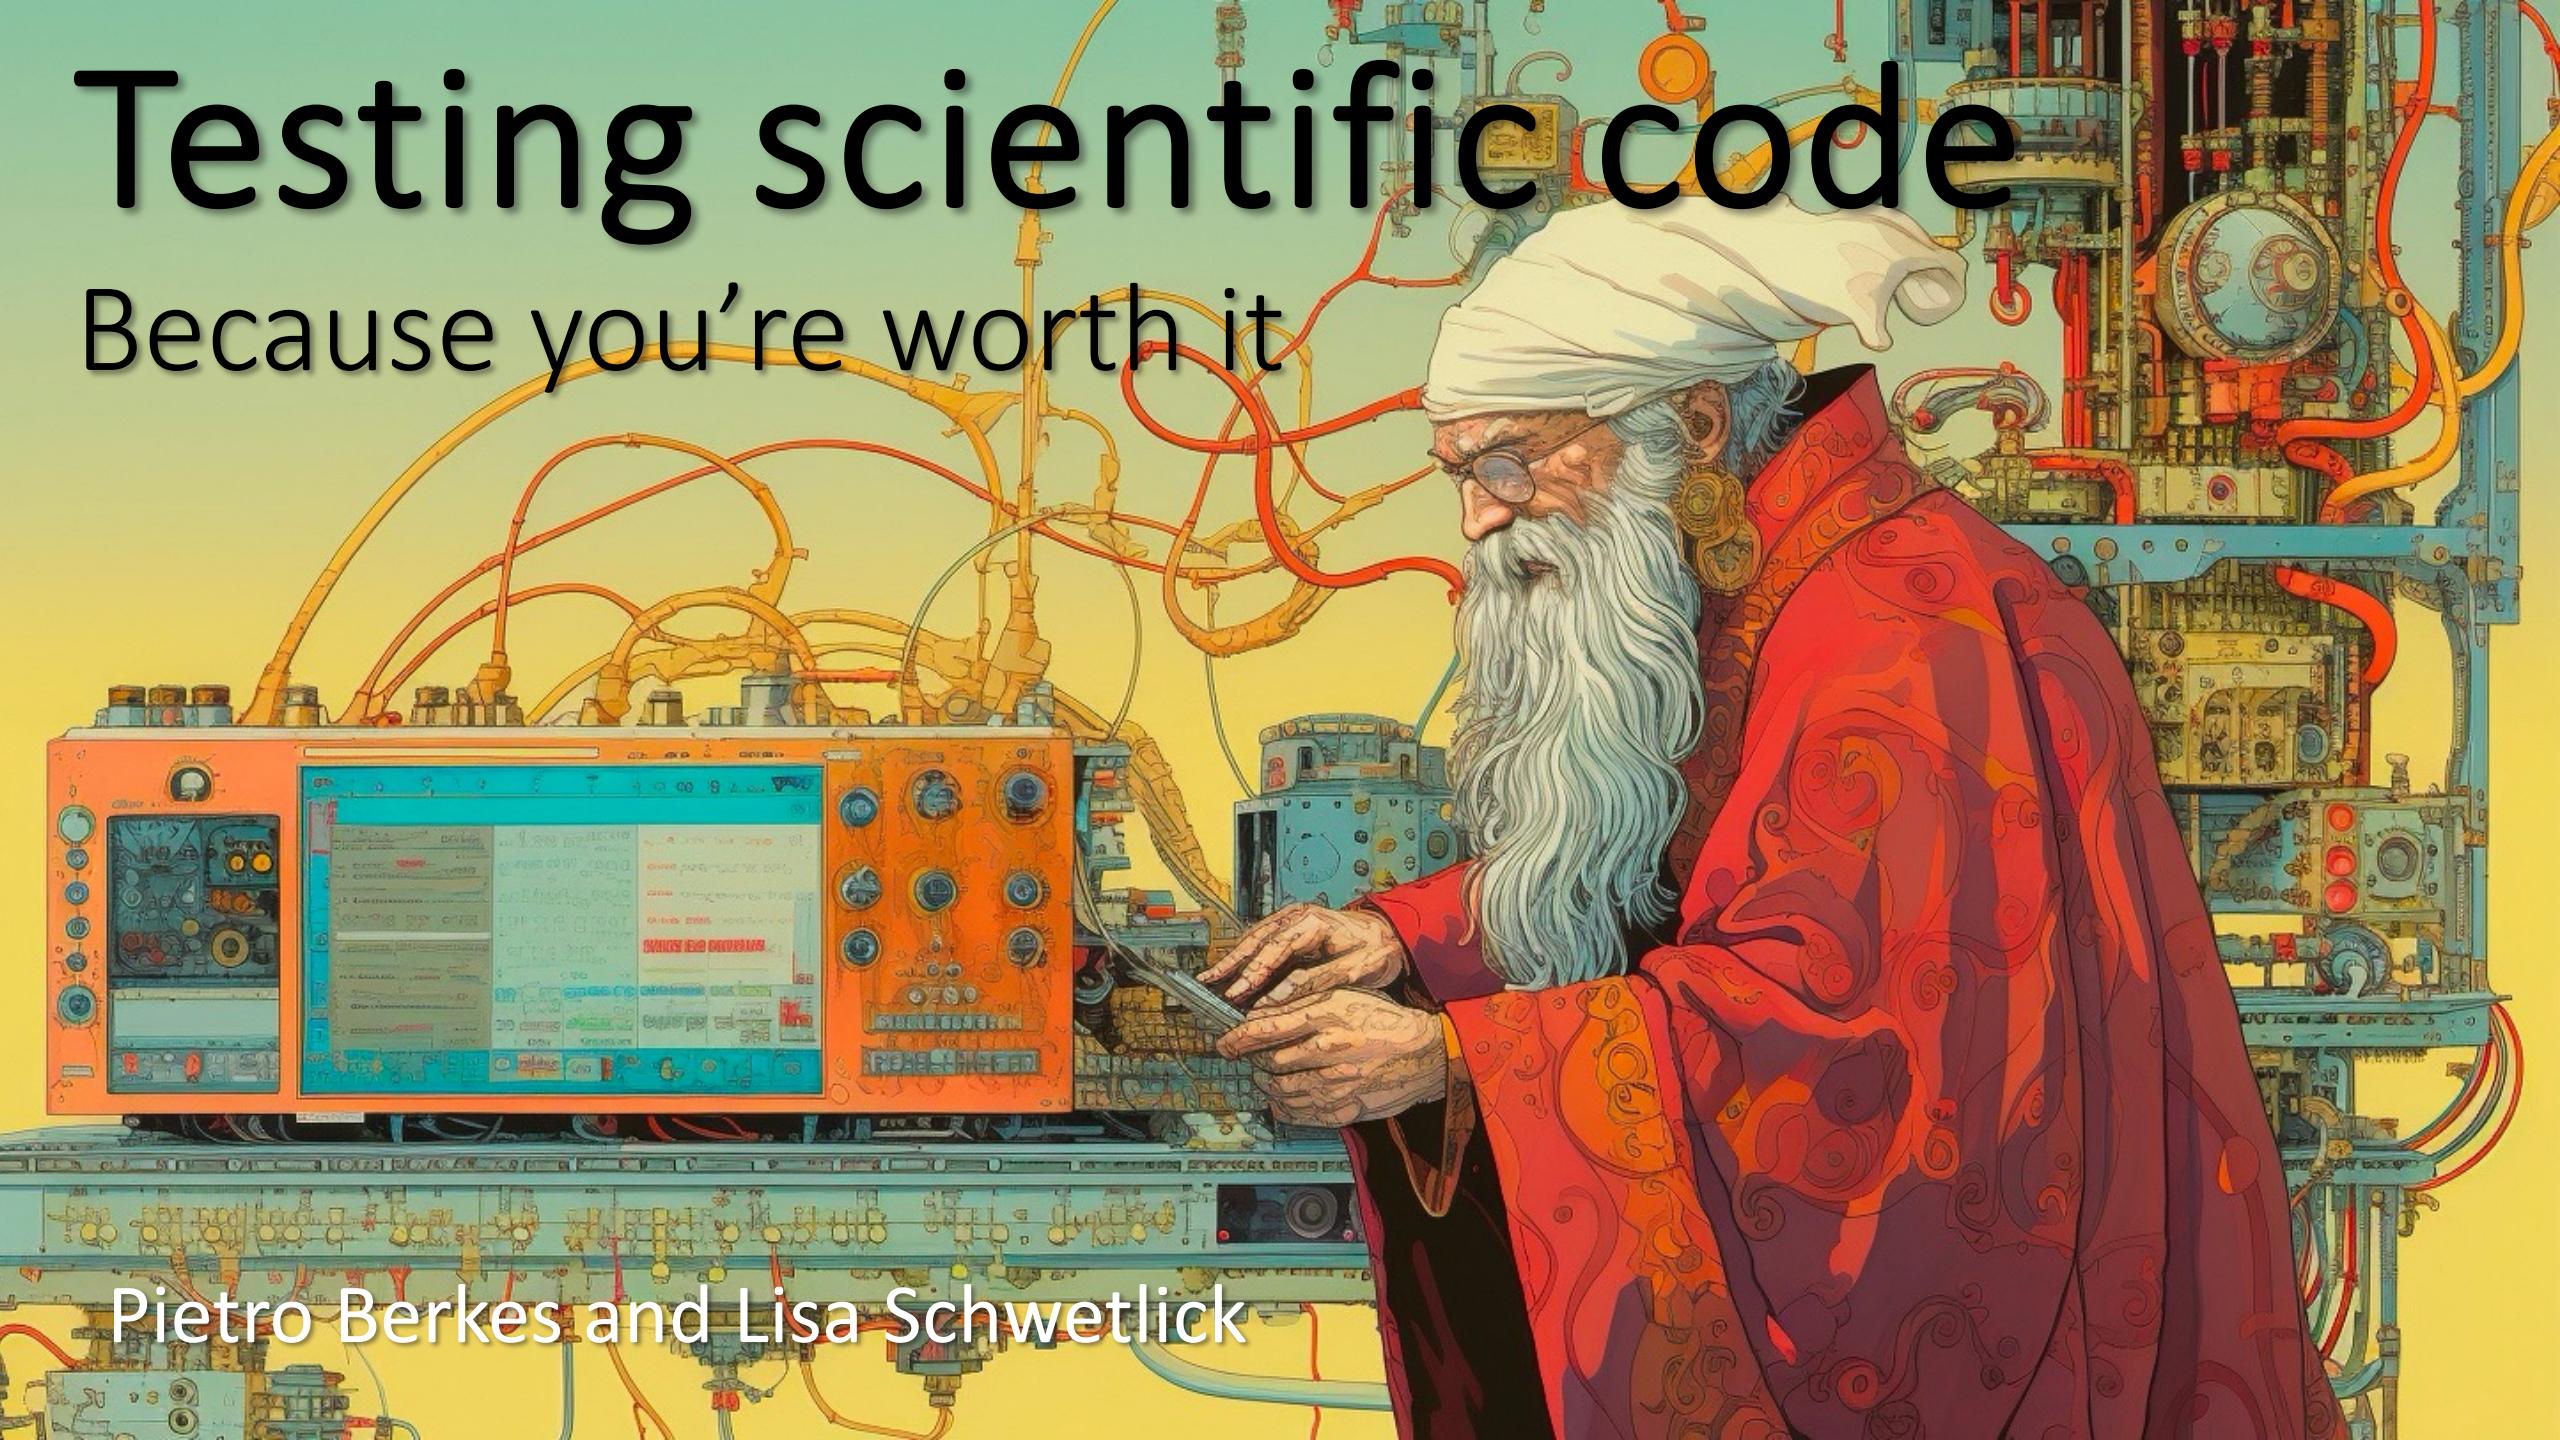


Testing scientific code

Because you're worth it



Pietro Berkes and Lisa Schwetlick



You, as the Master of Research

You start a new project and identify a number of possible leads.

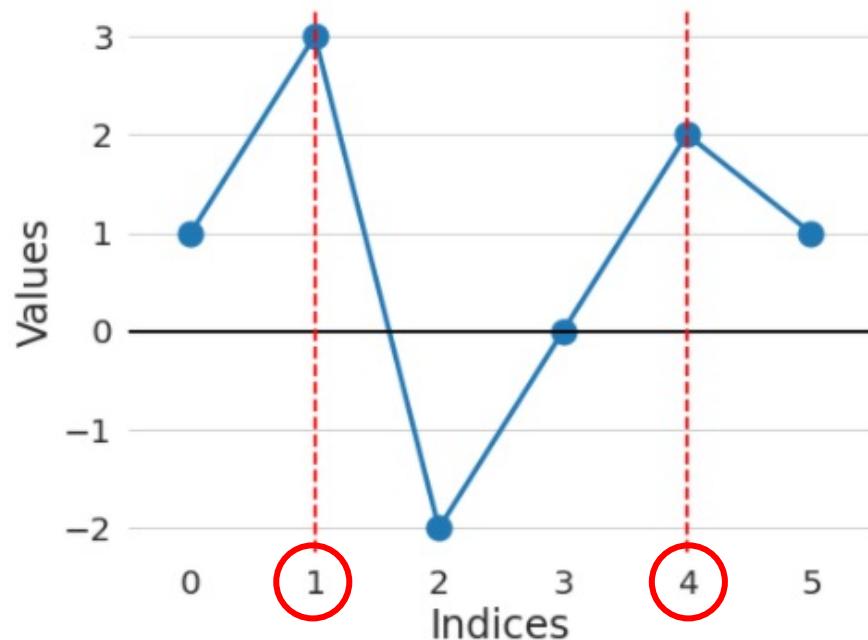
You **quickly develop a prototype** of the most promising ones; once a prototype is finished, you can **confidently decide** whether it is a dead end, or worth pursuing.

Once you find an idea on which it is worth spending energy, you take the prototype and **easily re-organize and optimize it** so that it scales up to the full size of your problem.

As expected, the scaled-up experiment delivers good results, and your next paper is under way.

Warm-up project

- Go to the directory called `hands_on/local_maxima`
- In the file called `local_maxima.py`, write a function `find_maxima` that finds the indices of local maxima in a list of numbers



For example,
`find_maxima([1, 3, -2, 0, 2, 1])`
should return
`[1, 4]`

Warm-up project

- Write a function `find_maxima` that finds the indices of local maxima in a list of numbers
- Check your solution with these inputs:
 - Input: [1, 3, -2, 0, 2, 1] Expected result: [1, 4]
 - Input: [4, 2, 1, 3, 1, 5] Expected result: [0, 3, 5]
 - Input: [] Expected result: []
 - Input: [1, 2, 2, 1] Expected result: [1] (or [2], or [1, 2])
 - Input: [1, 2, 2, 3, 1] Expected result: [3]

Testing basics

A test is just another function

- Imagine we wrote this new function, and we wanted to test it

```
def times_3(x):
    """Multiply x by 3.

    Parameters
    -----
    x : The item to multiply by 3.
    """
    return x * 3
```

Testing frameworks

- The collection of tests written to test a package is called a “test suite”
- Execution of a test suite is automated: external software runs the tests and provides reports and statistics
- Main testing frameworks for python:
 - unittest: in the standard library
 - **pytest: what is most commonly used**

```
===== test session starts =====
platform darwin -- Python 3.11.3, pytest-7.3.1, pluggy-1.0.0
collected 2 items

test_first.py::test_times_3_integer PASSED [ 50%]
test_first.py::test_times_3_string PASSED [100%]

===== 2 passed in 0.00s =====
```

Hands-on!

- Go to `hands_on/first`
 1. Discover all tests in all subdirectories
`pytest -v`
 2. Execute all tests in one module
`pytest -v test_first.py`
 3. Execute one single test
`pytest -v test_first.py::test_times_3_string`

Test suites in Python with pytest

- Writing tests with pytest is simple:
 - Tests are collected in files called `test_abc.py` , which usually contains tests for the functions defined in a corresponding module abc
 - Each test is a function called `test_jkl_feature`, and usually it tests feature feature of a function called jkl
 - Each test tests **one feature** in your code, and checks that it behaves correctly using “assertions”. An exception is raised if it does not work as expected.

Basic structure of test

- A good test is divided in three parts:
 - **Given:** Put your system in the right state for testing
 - Create data, initialize parameters, define constants...
 - **When:** Execute the feature that you are testing
 - Typically, one or two lines of code
 - **Then:** Compare outcomes with the expected ones
 - Define the expected result of the test
 - Set of *assertions* that check that the new state of your system matches your expectations

Assertions

- assert statements check that some condition is met, and raise an exception otherwise

- Check that statement is true/false:

```
assert 'Hi'.islower()      => fail  
assert not 'Hi'.islower() => pass
```

- Check that two objects are equal:

```
assert 2 + 1 == 3          => pass  
assert [2] + [1] == [2, 1] => pass  
assert 'a' + 'b' != 'ab'   => fail
```

- assert can be used to compare all sorts of objects, and pytest will take care of producing an appropriate error message

What a good test looks like

- What does a good test look like? What should I test?
- Good:
 - Short and quick to execute
 - Easy to read
 - Tests *one* thing
- Bad:
 - Relies on data files
 - Messes with “real-life” files, servers, databases

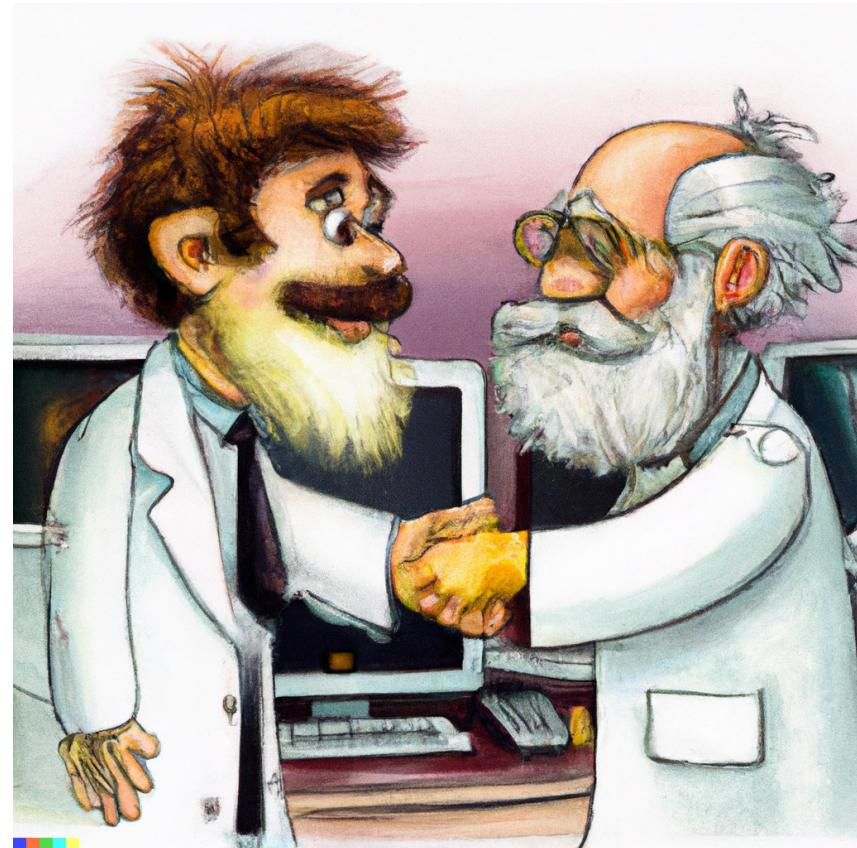
Write a working version of `find_maxima`, with testing

- **Read carefully the description of Issue #2 on GitHub**
- Submit a Pull Request for Issue #2
 - Fork the repository (if you haven't already)
 - Create a new branch on the fork called, e.g., `fix-2`
 - Solve the issue with one or more commits
 - Push the branch to your GitHub fork
 - On GitHub, go to “Pull Requests” and open a pull request against branch `main` of the official ASPP repository
 - In the PR description write “Fixes #2” somewhere, this is going to create an automatic link to the issue, and close the issue if the PR is merged

Testing is good for your self-esteem

- Immediately: Always be confident that your results are correct, whether your approach works or not
- In the future: **save your future self some trouble!**
- If you are left thinking “it’s cool but I cannot test *my* code because XYZ”, talk to us and we’ll show you how to do it ;-)

You, in 2024 → You, in 2025



Testing patterns

Floating point equality

- Real numbers are represented approximately as “floating point” numbers. When developing numerical code, we have to allow for approximation errors.
- Check that two numbers are approximately equal:

```
from math import isclose
def test_floating_point_math():
    assert isclose(1.1 + 2.2, 3.3)          => pass
```

- `abs_tol` controls the absolute tolerance:

```
assert isclose(1.121, 1.2, abs_tol=0.1)      => pass
assert isclose(1.121, 1.2, abs_tol=0.01)       => fail
```

- `rel_tol` controls the relative tolerance:

```
assert isclose(120.1, 121.4, rel_tol=0.1)     => pass
assert isclose(120.4, 121.4, rel_tol=0.01)      => fail
```

Testing with numpy arrays

```
def test_numpy_equality():
    x = np.array([1, 1])
    y = np.array([2, 2])
    z = np.array([3, 3])
    assert x + y == z
```

test_numpy_equality

```
def test_numpy_equality():
    x = numpy.array([1, 1])
    y = numpy.array([2, 2])
    z = numpy.array([3, 3])
>   assert x + y == z
E       ValueError: The truth value of an array with more than one element is ambiguous.
       Use a.any() or a.all()

code.py:47: ValueError
```

Testing with numpy arrays

- The module `np.testing` defines helper functions:

```
assert_equal(x, y)  
assert_allclose(x, y, rtol=1e-07, atol=0)
```

- If you need to check more complex conditions:

- `np.all(x)`: returns True if all elements of x are true
 - `np.any(x)`: returns True if any of the elements of x is true

- combine with `logical_and`, `logical_or`, `logical_not`:

```
# test that all elements of x are between 0 and 1  
assert all(logical_and(x > 0.0, x < 1.0))
```

Watch out for nans!

- In general, nan is not equal to itself (IEEE standard)

```
In [2]: np.nan == np.nan  
Out[2]: False
```

- `assert_equal` and `assert_allclose` consider nans equal by default

```
def test_allclose_with_nan():  
    x = np.array([1.1, np.nan])  
    y = np.array([2.2, np.nan])  
    z = np.array([3.3, np.nan])  
    assert_allclose(x + y, z)
```

```
test_numpy_equality.py::test_allclose_with_nan      PASSED
```

Common for-loop pattern for testing multiple cases

- Often these cases are collected in a single test:

```
def test_lower():
    # Given
    # Each test case is a tuple of (input, expected_result)
    test_cases = [('Hello wOrld', 'hello world'),
                  ('hi', 'hi'),
                  ('123 ([?', '123 ([?'),
                  ('', '')]

    for string, expected in test_cases:
        # When
        output = string.lower()
        # Then
        assert output == expected
```

The for-loop pattern can be improved

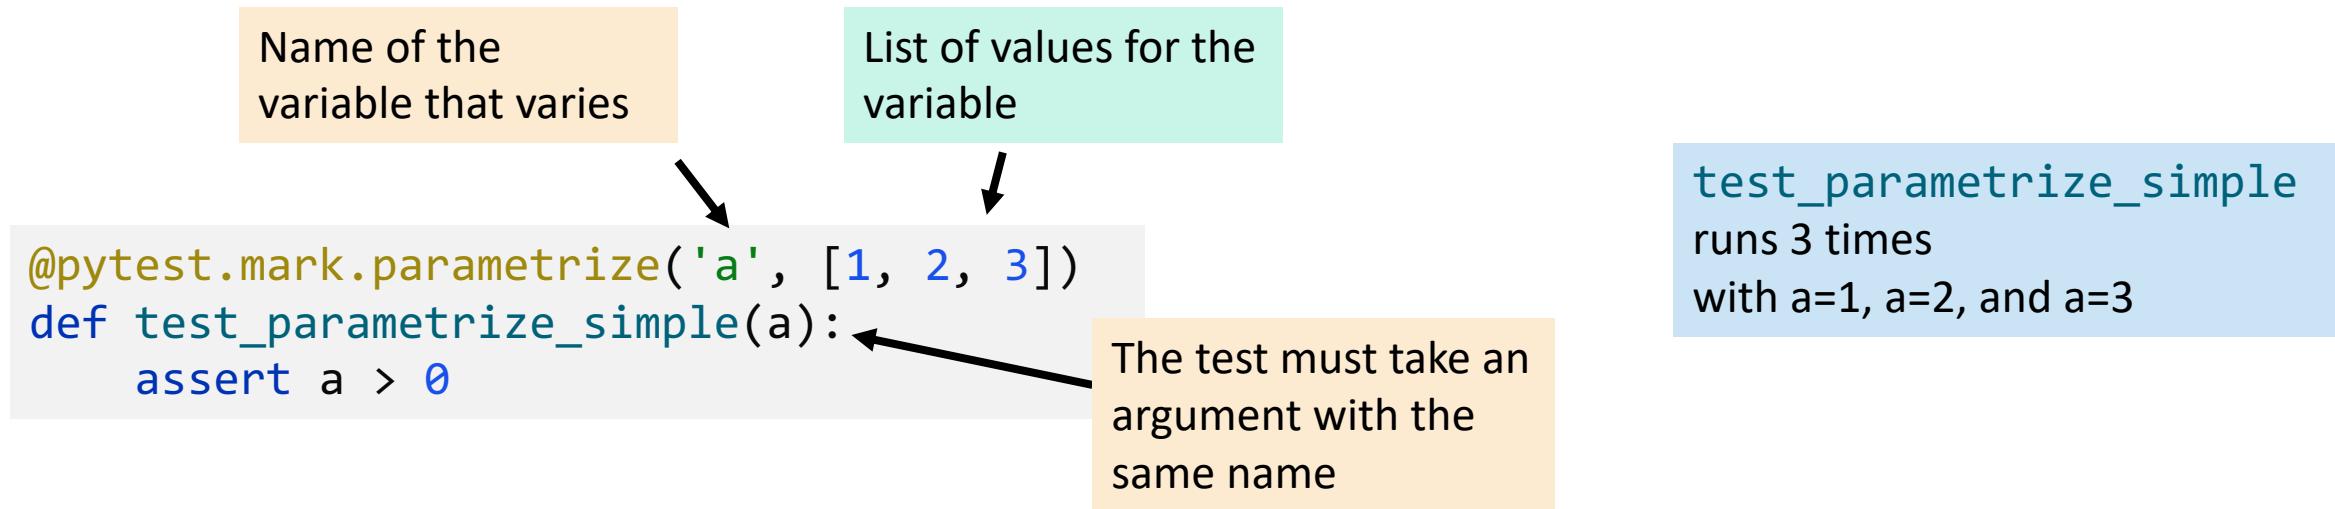
- It is repetitive to write the for-loop pattern
- If one of the cases break, it can be complicated to figure out which one
- pytest has many helpers for simplifying common testing cases!
- One of them is the `parametrize` decorator, that simplifies running the same test with multiple cases

Simple example

```
def test_for_loop_simple():
    cases = [1, 2, 3]
    for a in cases:
        assert a > 0
```

`test_for_loop_simple`
runs once and loops over
3 test cases

Simple example, with the `parametrize` decorator



Simple example, with the parametrize decorator

```
@pytest.mark.parametrize('a', [1, 2, 3])
def test_parametrize_simple(a):
    assert a > 0
```

Name of the variable that varies

List of values for the variable

The test must take an argument with the same name

```
===== test session starts =====
platform darwin -- Python 3.11.3, pytest-7.3.1, pluggy-1.0.0 -- /Users/pietro.berkes/miniconda3/envs/aspp/bin/python
cachedir: .pytest_cache
rootdir: /Users/pietro.berkes/o/ASPP/testing_project/demos
plugins: anyio-3.5.0
collected 3 items

test_parametrize.py::test_parametrize_simple[1] PASSED
[ 33%]
test_parametrize.py::test_parametrize_simple[2] PASSED
[ 66%]
test_parametrize.py::test_parametrize_simple[3] PASSED
[100%]

===== 3 passed in 0.00s =====
```

pytest automatically creates one separate test for each test case

Example with multiple values

- This is a more typical case with several input values and the expected result of the test

```
def test_for_loop_multiple():
    cases = [
        (1, 'hi', 'hi'),
        (2, 'no', 'nono')
    ]
    for a, b, expected in cases:
        result = b * a
        assert result == expected
```

test_for_loop_multiple
runs once and loops over
2 test cases

Same example, with the parametrize decorator

Name of all the variables,
separated by commas in
one string

List of tuples with the
values for each variable,
one for each test case

```
@pytest.mark.parametrize('a, b, expected', [(1, 'hi', 'hi'), (2, 'no', 'nono')])  
def test_parametrize_multiple(a, b, expected):  
    result = b * a  
    assert result == expected
```

The test must take
arguments with the
same names as in the
string

test_parametrize_multiple
runs 2 times with
1) a=1 b='hi' expected='hi'
and
2) a=2 b='no', expected='nono'

Same example, with the parametrize decorator

Name of all the variables,
separated by commas in
one string

List of tuples with the
values for each variable,
one for each test case

```
@pytest.mark.parametrize('a, b, expected', [(1, 'hi', 'hi'), (2, 'no', 'nono')])
def test_parametrize_multiple(a, b, expected):
    result = b * a
    assert result == expected
```

The test must take
arguments with the
same names as in the
string

```
[4] pytest -v test_parametrize.py::test_parametrize_multiple
===== test session starts =====
platform darwin -- Python 3.11.3, pytest-7.3.1, pluggy-1.0.0 -- /Users/pietro.berkes/miniconda3/envs/aspp/bin/python
cachedir: .pytest_cache
rootdir: /Users/pietro.berkes/o/ASPP/testing_project/demos
plugins: anyio-3.5.0
collected 2 items

test_parametrize.py::test_parametrize_multiple[1-hi-hi] PASSED [ 50%]
test_parametrize.py::test_parametrize_multiple[2-no-nono] PASSED [100%]

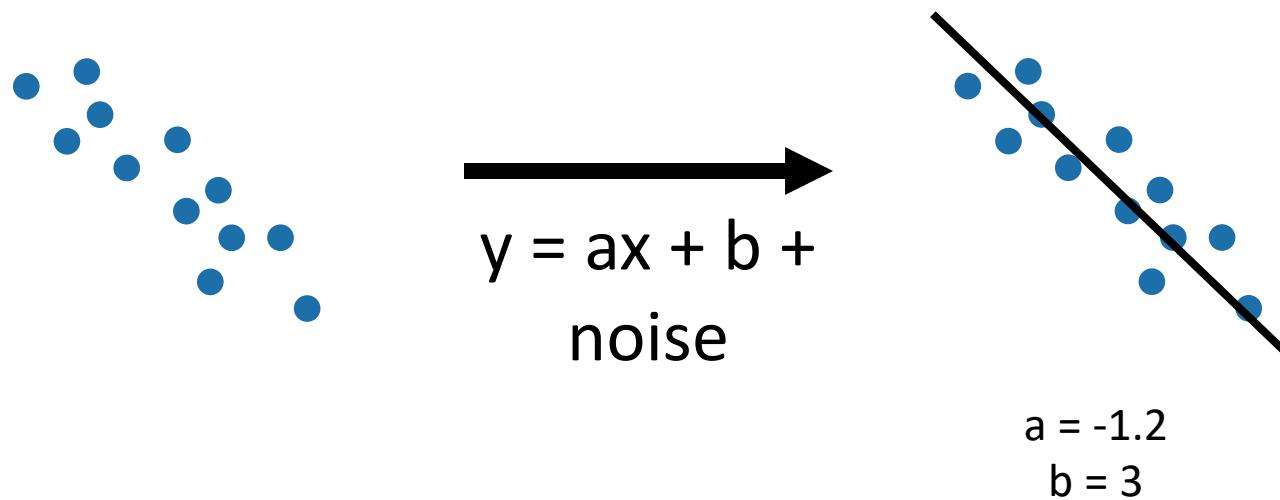
===== 2 passed in 0.01s =====
```

pytest automatically
creates one separate
test for each test case

Strategies for testing learning algorithms

- Learning algorithms can get stuck in local maxima, the solution for general cases might not be known (e.g., unsupervised learning)
- Turn your validation cases into tests
- Stability tests:
 - Start from final solution; verify that the algorithm stays there
 - Start from solution and add a small amount of noise to the parameters; verify that the algorithm converges back to the solution
- Parameter Recovery: Generate synthetic data from the model with known parameters, then test that the code can learn the parameters back

Learning algorithms fit the parameters of a model to observed data



Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters

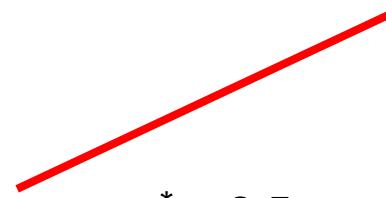


$$a^* = 0.5$$

$$b^* = -1.3$$

Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters



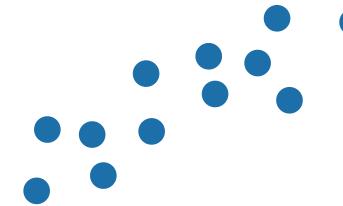
$$a^* = 0.5$$
$$b^* = -1.3$$



$$y = a^* x + b^*$$

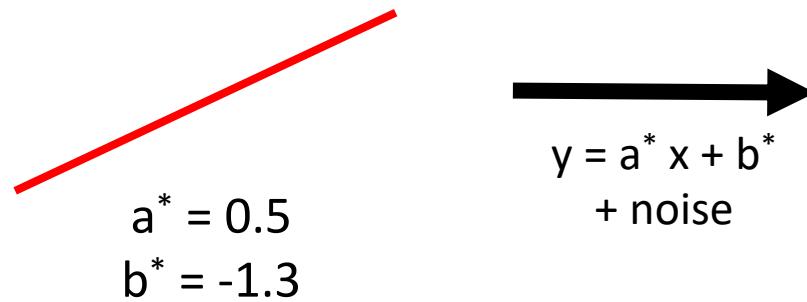
+ noise

2) Generate synthetic data

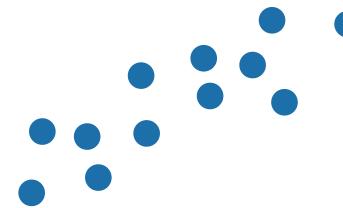


Generate synthetic data from the model to test the learning algorithm by recovering the parameters

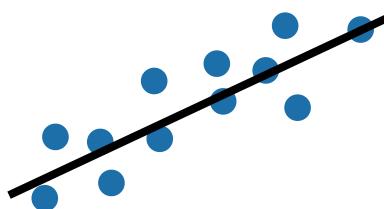
1) Fix initial parameters



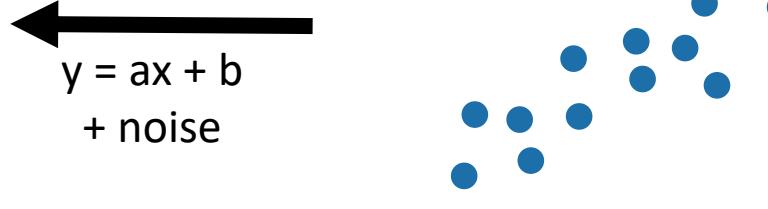
2) Generate synthetic data



$$\begin{aligned} a &= 0.5098 \\ b &= -1.287 \end{aligned}$$

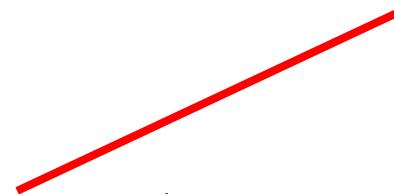


3) Run the algorithm

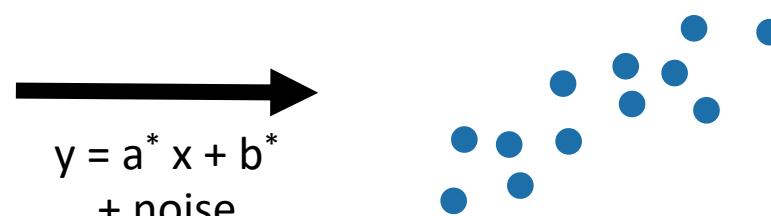


Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters


$$a^* = 0.5$$
$$b^* = -1.3$$

2) Generate synthetic data

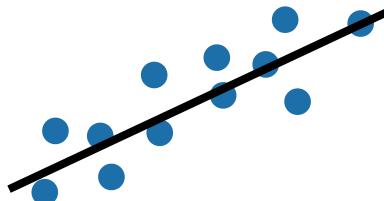


4) Compare



$$a = 0.5098$$
$$b = -1.287$$

3) Run the algorithm



Randomness in Testing

- Using randomness in testing can be useful
 - To check that the code is stable and works correctly in many different cases
 - To find corner cases or numerical problems

```
def test_logistic_fit_randomized():
    random_state = np.random.RandomState(SEED)
    for _ in range(100):
        x0 = random_state.uniform(0.0001, 0.9999)
        r = round(random_state.uniform(0.001, 3.999), 3)

        xs = iterate_f(it=17, x0=x0, r=r)
        recovered_r = fit_r(xs)

    assert_allclose(r, recovered_r, atol=1e-3)
```



Random Seeds and Reproducibility

- When running tests that involve randomness and some test doesn't pass it is vital to be able to reproduce that test exactly!
- Computers produce pseudo-random numbers: setting a seed resets the basis for the random number generator
- This is essential for reproducibility
- At a minimum, you should manually set the seed for each of your random tests

```
SEED = 42
random_state = np.random.RandomState(SEED)
random_state.rand()
```

A Pytest Solution

- Non-scientific coding uses random testing more rarely, so there is no helper tools for that in pytest
- However, in scientific coding it is quite common
- What do we want?
 - For each (random) test there should be a seed
 - For each run of the test, the seed should be different
 - That seed should be printed with the test result
 - It needs to be possible to explicitly run the test again with that seed!

Fixtures (minimal solution)

- Fixtures are functions that are run before the tests are executed

```
import numpy as np
import pytest

# set the random seed for once here
SEED = np.random.randint(0, 2**31)

@pytest.fixture
def random_state():
    print(f'Using seed {SEED}')
    random_state = np.random.RandomState(SEED)
    return random_state

def test_something(random_state):
    random_state.rand()
```

If an input argument of a test matches the name of a fixture, then the fixture is called and the return value assigned to the argument.

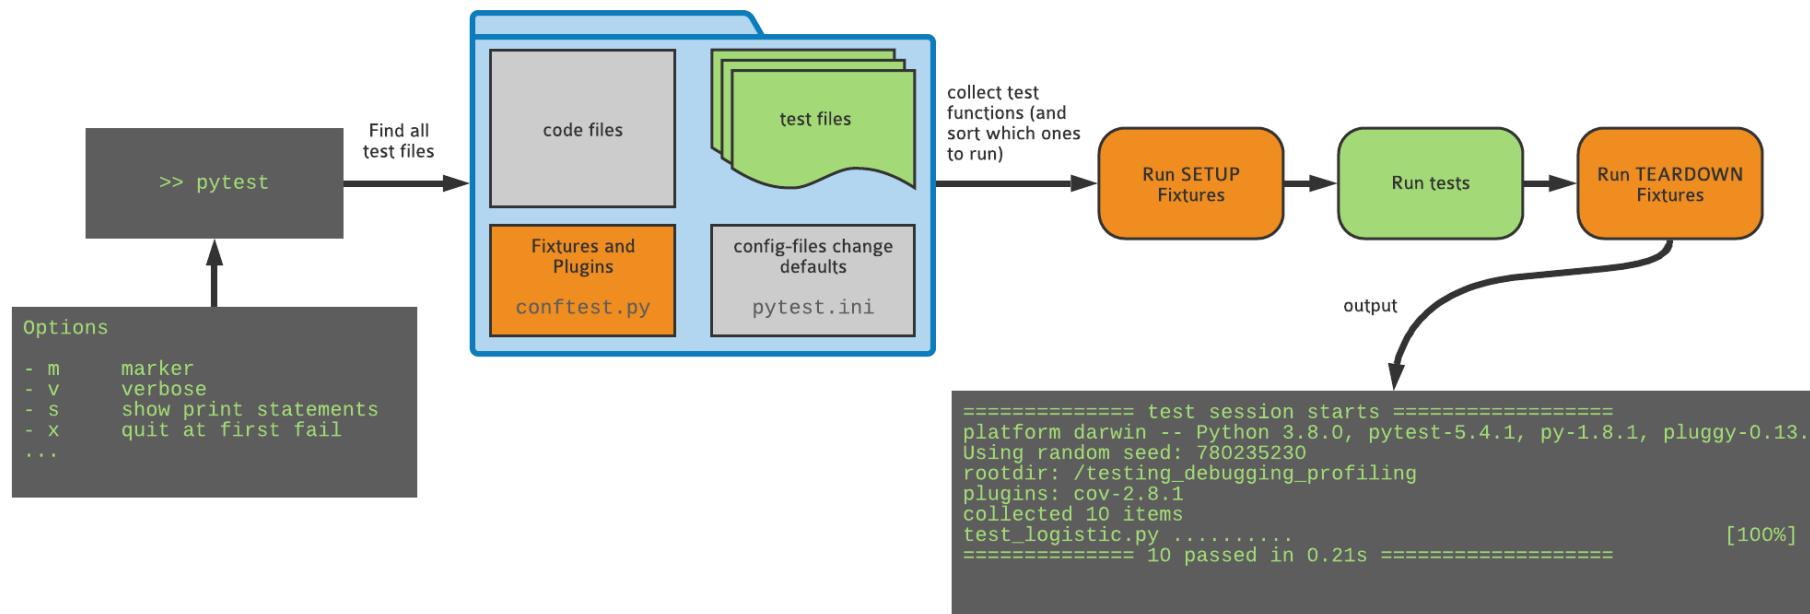
pytest handles that automatically as part of running the test suite

Hands On!

- a) Write a randomized test that checks that `fit_r` can recover r for any random value of x_0 and r
- b) Add a fixture at the top of your test file, that lets you print the seed to the console.

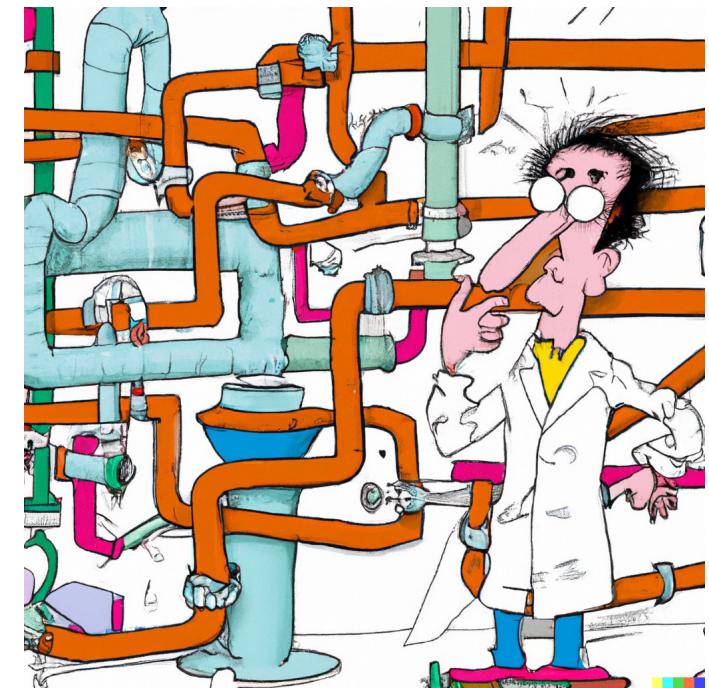
```
[└$ pytest
               test session starts
platform darwin -- Python 3.8.0, pytest-5.4.1, py-1.8.1, pluggy-0.13.1
Using random seed: 892358865]
```

What happens when you run pytest



Fixtures (real solution)

- `conftest.py` is a special pytest config file (don't import it!)
- `conftest.py` can be used to define custom behavior or plugins. Fixtures can also be defined here, so that they can be used by all tests.
- See the file `hands_on/randomness/_conftest.py` in the repo you forked. If you move it to the main folder and rename it, the functions defined there select a seed for each test and allow you to pass a seed on the command line using `--seed 123`



Decorating “special” tests

- `@xfail`: Expected failure, outputs an “x” (or “X”) in the report

```
@pytest.mark.xfail  
def test_something():  
    ...
```

- `@skip`: Skip test, useful e.g. when the feature doesn’t exist yet

```
@pytest.mark.skip(reason="functionality not yet  
implemented")  
def test_something():  
    ...
```

- `@skipif`: Skip the test if a condition is met, useful for tests that only works on a specific platform, or for a specific version of Python

```
@pytest.mark.skipif(sys.version_info < (3, 10),  
                    reason="requires python3.10 or higher")  
def test_something():  
    ...
```

Marking tests with custom markers

- If you have lots of tests, you can categorize them with your own markers
 - although for custom mark names you need to register the marks “pytest.ini”
 - <https://docs.pytest.org/en/7.1.x/example/markers.html#registering-markers>
- Example:
 - Smoke tests check for really basic features: run these frequently
 - Other tests may be many or too slow to run every time and test for more edge cases

```
@pytest.mark.smoke
def test_something_basic():
    ...
```

```
> pytest -m smoke
> pytest -m "smoke and not slow"
```

Writing temporary files: `tmp_path`

- To test functions that write to disk without leaving around the files when the test is finished, use the `tmp_path` fixture
- The value of `tmp_path` is a `pathlib.Path` object
- The directory is created at the start of the test, and removed at the end

```
def test_create_file(tmp_path):  
    d = tmp_path / "sub"  
    d.mkdir()  
    p = d / "hello.txt"  
    content = "some random text"  
    p.write_text(content)  
    assert p.read_text() == content  
    assert len(list(tmp_path.iterdir())) == 1
```

All you need to do is add an argument with this exact name